

# rice-disease-resnet50

May 7, 2024

```
[6]: import os
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, models, transforms
from torch.utils.data import DataLoader
```

```
[7]: # Define device
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

# Define data directory paths
data_dir = "rice_diseases"
train_dir = os.path.join(data_dir, "train")
val_dir = os.path.join(data_dir, "val")

# Define data transforms
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}
```

```
[8]: # Create datasets
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
    ↪data_transforms[x]) for x in ['train', 'val']}

# Create data loaders
```

```

dataloaders = {x: DataLoader(image_datasets[x], batch_size=32, shuffle=True,
    ↪num_workers=4) for x in ['train', 'val']}

# Define the model architecture (using ResNet-50)
model = models.resnet50(pretrained=True)
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, len(image_datasets['train'].classes))

# Move model to device
model = model.to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

```

c:\Users\fr5009tu\AppData\Local\Programs\Python\Python311\Lib\site-packages\torchvision\models\\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

warnings.warn(

c:\Users\fr5009tu\AppData\Local\Programs\Python\Python311\Lib\site-packages\torchvision\models\\_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet50\_Weights.IMAGENET1K\_V1`. You can also use `weights=ResNet50\_Weights.DEFAULT` to get the most up-to-date weights.

warnings.warn(msg)

```

[9]: # Train the model
def train_model(model, criterion, optimizer, num_epochs):
    for epoch in range(num_epochs):
        print(f'Epoch {epoch+1}/{num_epochs}')
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

            running_loss = 0.0
            running_corrects = 0

            # Iterate over data.
            for inputs, labels in dataloaders[phase]:
                inputs = inputs.to(device)

```

```

labels = labels.to(device)

# Zero the parameter gradients
optimizer.zero_grad()

# Forward pass
with torch.set_grad_enabled(phase == 'train'):
    outputs = model(inputs)
    _, preds = torch.max(outputs, 1)
    loss = criterion(outputs, labels)

# Backward + optimize only if in training phase
if phase == 'train':
    loss.backward()
    optimizer.step()

# Statistics
running_loss += loss.item() * inputs.size(0)
running_corrects += torch.sum(preds == labels.data)

epoch_loss = running_loss / len(image_datasets[phase])
epoch_acc = running_corrects.double() / len(image_datasets[phase])

print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

```

```

[10]: # Train the model
train_model(model, criterion, optimizer, num_epochs=10)

```

```

Epoch 1/10
-----
train Loss: 0.8500 Acc: 0.5920
val Loss: 0.9071 Acc: 0.6667
Epoch 2/10
-----
train Loss: 0.5180 Acc: 0.7713
val Loss: 0.5162 Acc: 0.7793
Epoch 3/10
-----
train Loss: 0.4154 Acc: 0.8201
val Loss: 0.3825 Acc: 0.8161
Epoch 4/10
-----
train Loss: 0.4130 Acc: 0.8098
val Loss: 0.3478 Acc: 0.8437
Epoch 5/10
-----
train Loss: 0.3464 Acc: 0.8402

```

```

val Loss: 0.2900 Acc: 0.8713
Epoch 6/10
-----
train Loss: 0.3427 Acc: 0.8397
val Loss: 0.3329 Acc: 0.8414
Epoch 7/10
-----
train Loss: 0.3083 Acc: 0.8552
val Loss: 0.2331 Acc: 0.9034
Epoch 8/10
-----
train Loss: 0.3229 Acc: 0.8511
val Loss: 0.2488 Acc: 0.8805
Epoch 9/10
-----
train Loss: 0.3008 Acc: 0.8718
val Loss: 0.3104 Acc: 0.8598
Epoch 10/10
-----
train Loss: 0.2965 Acc: 0.8626
val Loss: 0.2231 Acc: 0.8966

```

```

[11]: import pickle

      # Save the entire model object
      with open('rice_disease_model.pkl', 'wb') as f:
          pickle.dump(model, f)

```

```

[12]: # Step 1: Load the model from pickle
      with open('rice_disease_model.pkl', 'rb') as f:
          new_model = pickle.load(f)

```

```

[13]: # Train the model
      train_model(new_model, criterion, optimizer, num_epochs=10)

```

```

Epoch 1/10
-----
train Loss: 0.2880 Acc: 0.8632
val Loss: 0.2130 Acc: 0.8897
Epoch 2/10
-----
train Loss: 0.2900 Acc: 0.8644
val Loss: 0.2198 Acc: 0.8851
Epoch 3/10
-----
train Loss: 0.2908 Acc: 0.8770
val Loss: 0.2100 Acc: 0.8897
Epoch 4/10

```

```

-----
train Loss: 0.2650 Acc: 0.8741
val Loss: 0.2118 Acc: 0.8851
Epoch 5/10
-----
train Loss: 0.2995 Acc: 0.8609
val Loss: 0.2094 Acc: 0.8851
Epoch 6/10
-----
train Loss: 0.2951 Acc: 0.8678
val Loss: 0.2048 Acc: 0.8897
Epoch 7/10
-----
train Loss: 0.2876 Acc: 0.8747
val Loss: 0.2127 Acc: 0.8874
Epoch 8/10
-----
train Loss: 0.3078 Acc: 0.8546
val Loss: 0.2053 Acc: 0.8897
Epoch 9/10
-----
train Loss: 0.2956 Acc: 0.8603
val Loss: 0.2044 Acc: 0.8851
Epoch 10/10
-----
train Loss: 0.3032 Acc: 0.8690
val Loss: 0.2173 Acc: 0.8943

```

```

[14]: # Save the entire model object
with open('rice_disease_model.pkl', 'wb') as f:
    pickle.dump(new_model, f)

```

```

[15]: # Step 1: Load the model from pickle
with open('rice_disease_model.pkl', 'rb') as f:
    new_model = pickle.load(f)

```

```

[16]: # Train the model
train_model(new_model, criterion, optimizer, num_epochs=10)

```

```

Epoch 1/10
-----
train Loss: 0.2894 Acc: 0.8770
val Loss: 0.2091 Acc: 0.8874
Epoch 2/10
-----
train Loss: 0.2736 Acc: 0.8810
val Loss: 0.2096 Acc: 0.8920
Epoch 3/10

```

```

-----
train Loss: 0.2726 Acc: 0.8810
val Loss: 0.2122 Acc: 0.8874
Epoch 4/10
-----
train Loss: 0.3106 Acc: 0.8609
val Loss: 0.2081 Acc: 0.8897
Epoch 5/10
-----
train Loss: 0.2931 Acc: 0.8667
val Loss: 0.2013 Acc: 0.8897
Epoch 6/10
-----
train Loss: 0.2706 Acc: 0.8764
val Loss: 0.2127 Acc: 0.8897
Epoch 7/10
-----
train Loss: 0.2879 Acc: 0.8701
val Loss: 0.2165 Acc: 0.8805
Epoch 8/10
-----
train Loss: 0.2919 Acc: 0.8730
val Loss: 0.2161 Acc: 0.8897
Epoch 9/10
-----
train Loss: 0.2761 Acc: 0.8718
val Loss: 0.2024 Acc: 0.8897
Epoch 10/10
-----
train Loss: 0.2899 Acc: 0.8678
val Loss: 0.2069 Acc: 0.8897

```

```

[17]: # Save the entire model object
with open('rice_disease_model.pkl', 'wb') as f:
    pickle.dump(new_model, f)

```

```

[18]: # Step 1: Load the model from pickle
with open('rice_disease_model.pkl', 'rb') as f:
    new_model = pickle.load(f)

```

```

[19]: # Train the model
train_model(new_model, criterion, optimizer, num_epochs=10)

```

```

Epoch 1/10
-----
train Loss: 0.2970 Acc: 0.8603
val Loss: 0.2102 Acc: 0.8920
Epoch 2/10

```

```

-----
train Loss: 0.3027 Acc: 0.8580
val Loss: 0.2216 Acc: 0.8943
Epoch 3/10
-----
train Loss: 0.2952 Acc: 0.8655
val Loss: 0.2078 Acc: 0.8920
Epoch 4/10
-----
train Loss: 0.2989 Acc: 0.8632
val Loss: 0.2142 Acc: 0.8920
Epoch 5/10
-----
train Loss: 0.2939 Acc: 0.8649
val Loss: 0.2115 Acc: 0.8897
Epoch 6/10
-----
train Loss: 0.2983 Acc: 0.8724
val Loss: 0.2084 Acc: 0.8989
Epoch 7/10
-----
train Loss: 0.2948 Acc: 0.8707
val Loss: 0.2134 Acc: 0.8920
Epoch 8/10
-----
train Loss: 0.2694 Acc: 0.8805
val Loss: 0.2148 Acc: 0.8851
Epoch 9/10
-----
train Loss: 0.2865 Acc: 0.8649
val Loss: 0.2074 Acc: 0.8943
Epoch 10/10
-----
train Loss: 0.2949 Acc: 0.8695
val Loss: 0.2104 Acc: 0.8874

```

```

[20]: # Save the entire model object
with open('rice_disease_model.pkl', 'wb') as f:
    pickle.dump(new_model, f)

```